




S&DS 265 / 565
Introductory Machine Learning


Random Forests

February 11


Yale

Friendly reminders



- Assn 2 due next week
- Quiz 2 live after class today
- Midterm March 4, in class
- Practice midterms (and solutions) posted on Canvas
- Review sessions will be organized



X_1



X_1



Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

- Both produce simple predictions (averages/maximally occurring) based on “neighborhoods” in the predictor space.

Trees vs. other methods

Decision trees are similar in spirit to k -nearest neighbors.

- Both produce simple predictions (averages/maximally occurring) based on “neighborhoods” in the predictor space.
- Neighborhoods chosen very differently

Trees vs. other methods

Recall that linear regression fits models of the form

$$f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$$

Trees vs. other methods

Recall that linear regression fits models of the form





$$f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$$

Regression trees are like fitting linear regression models with a bunch of indicators

$$f(X) = \sum_{j=1}^J \beta_j \mathbb{1} \{X \in R_j\}$$

Trees vs. other methods


Are trees always better than linear methods?



Summary so far

- Trees give interpretable, nonlinear prediction rules
- Deep trees have low bias, high variance
- Shallow trees have high bias, low variance
- Deep trees are pruned back using cross-validation to find best bias/variance tradeoff.

Random Forests




Suppose that $Y_1, Y_2, \dots, Y_n \sim N(\theta, \sigma^2)$. Let $\bar{Y}_n = \frac{1}{n} \sum Y_i$ be the average. Then

$$\mathbb{E}(\bar{Y}_n) = \theta$$

If Y_i are uncorrelated then


$$\text{Var}(\bar{Y}_n) = \frac{\sigma^2}{n}$$

Random Forests



- Grow many trees and average their predictions
- Trees are grown deep, to have low bias, but high variance
- To “decorrelate” the trees and reduce variance, introduce randomness

Random Forests



- Grow many trees and average their predictions
- Trees are grown deep, to have low bias, but high variance
- To “decorrelate” the trees and reduce variance, each tree is
 - ▶ grown on a bootstrap sample of the data
 - ▶ grown with random subsets of the predictors at each split
- Tree growing can be done in parallel

Leo Breiman—"Keep it simple"



Random Forests Algorithm

- ① For $b = 1$ to B :
 - (a) Draw a bootstrap sample Z^* of size n from the training data
 - (b) Grow a random-forest tree T_b to the bootstrapped data, recursively repeating following steps, until minimum node size reached:
 - i. Select m variables at random from the p variables
 - ii. Pick the best variable/split-point among the m
 - iii. Split the node into two children nodes
- ② Output the ensemble of trees $\{T_b\}_{b=1}^B$.

Random Forests Algorithm

To make a prediction at a new point x :

Regression: Average $\hat{f}_{rf}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$

Classification: Majority vote of the individual trees


Out of bag (OOB) prediction

A bootstrap sample, or “bag” of data, is a set of n data points sampled with replacement from the original set of n data points. Will contain repetitions.

Each tree is grown on such a sample, which contains about $\frac{2}{3}$ of the original data (with repetitions). The remaining $\frac{1}{3}$ can be used as validation data

- For each observation $z_i = (x_i, y_i)$, construct its random forest predictor by averaging only those trees corresponding to bootstrap samples in which z_i did not appear
- Thus, cross-validation can be performed “along the way”

Out of bag (OOB) prediction



More detail: Bootstrap samples

- We have n data points
- Randomly sample exactly n of them *with replacement*
- Contains about $\frac{2}{3}$ of the data, with duplicates
- This is a “bag” of n data points

Chance a particular example does not appear in sample:

$$\left(1 - \frac{1}{n}\right)^n \rightarrow \frac{1}{e} \approx 0.37$$

More detail: Out-of-Bag error estimation

Each bagged tree uses about $2/3$ of all observations (with repeats).

The remaining data—*out-of-bag* (OOB) observations—can be put to good use.

More detail: Out-of-Bag error estimation

- Consider bags $\mathcal{B}_i \subset \{1, 2, \dots, B\}$ where (x_i, y_i) does *not* appear
- We average over those trees to form a prediction:

$$\hat{y}_i = \frac{1}{|\mathcal{B}_i|} \sum_{b \in \mathcal{B}_i} T_b(x_i)$$

- OOB error estimate is then $\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2$

Performance on email spam task

Single tree: 8.7%

Random forest: 5.1%

(standard error of the estimates is $\approx 0.6\%$)

Variable importance

Random forests improve upon the predictive ability of trees, but this sacrifices the interpretability of the single tree model

Variable importance

Random forests improve upon the predictive ability of trees, but this sacrifices the interpretability of the single tree model

A good tool for interpreting a forest is **variable importance**


Variable importance

Random forests improve upon the predictive ability of trees, but this sacrifices the interpretability of the single tree model


A good tool for interpreting a forest is **variable importance**

Variable importance is the amount that the RSS (or other loss) is reduced due to splits over a given predictor, averaged over all trees in the forest


Variable importance



Variable importance



Random forest MSE



Understanding the bias-variance tradeoff for random forests

- Each tree has low (squared) bias, because it is deep
- The bootstrap sample and random subset of questions allowed at a given split result in “diversity” among of the trees
- This diversity translates to decorrelation, and tends to reduce the variance
- As we increase the number m in the random subset of questions allowed at each split, the complexity increases—variance goes up, squared bias goes down

Please step through the notebook

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [2]: titanic_train = pd.read_csv('https://raw.githubusercontent.com/minsuk-heo/kaggle-titanic/master/input/train.csv')
titanic_test = pd.read_csv('https://raw.githubusercontent.com/minsuk-heo/kaggle-titanic/master/input/test.csv')
titanic_train
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cummings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...	...	...	...	...	...	...	...	...	...	...	...	...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q

891 rows × 12 columns

Summary: Trees and Forests

- Trees give interpretable, nonlinear prediction rules
- Deep trees have low bias, high variance
- Random forests are a way of combining trees
- Want: Different trees should capture different aspects of the data
- How: Grow each tree on a random (bootstrap) sample of the data and choosing from a random set of questions at each split