

Statistics and Data Science 265

Introductory Machine Learning

January 14

Yale

Outline

- Administrative items
- Elements of Python
- Demos

Office Hours

OH will be posted on Canvas and announced to the class.

Recordings

- Lectures will be recorded
 - Available on Canvas under “Media”
 - Please attend class and only use recordings for review
-

Upcoming

- Assignment 1 will be posted next week
- Topics: Classification and regression

Recap

- Last time: Course overview, logistics
- Any questions?

Plan for Today

- Python elements
- Pandas and linear regression example
- Basics of classification, regression, overfitting

Python primer: Concepts

- Python types: lists, tuples, strings, dictionaries
- Basics of iteration
- Comprehensions
- Arithmetic
- Printing
- NumPy and multi-dimensional arrays
- Array math
- DataFrames and pandas
- Matplotlib and basic plotting

Python elements

+ Code + Text

Python and Jupyter essentials for iML

This notebook was adapted from multiple resources including the Data8 curriculum, [Yale EENG201](#), and [Stanford CS231](#). It is intended to give you a quick "jumpstart" and introduction to the tools that we will use throughout the course, based on Python, Jupyter notebooks, and essential useful packages like `numpy` and `pandas`.

It's important to recognize that practice is crucial here—you need to write code and implement things, making mistakes along the way, to gain proficiency in this material.

 Subtopics marked with the scream icon are a little more advanced, and can be skipped on a first reading.

Get Started

Different ways to run Python

1. Create a file using editor, then: `$ python myscript.py`
2. Run interpreter interactively `$ python`
3. Use a Python environment, e.g. Anaconda or Google Colab

We recommend Anaconda:

- easy to install
- easy to add additional packages
- allows creation of custom environments

But Google Colab is also a good option. We plan to create a video on how to use Google Colab.

Resources

- Anaconda Python: <https://www.continuum.io>
- Jupyter notebooks: jupyter-notebook.readthedocs.io
- PyCharm debugger: www.jetbrains.com
- *Introducing Python*, Bill Lubanovic, O'Reilly
- *Python in a Nutshell*, Alex Martelli et al., O'Reilly
- *Python Cookbook*, David Beazley, Brian K. Jones, O'Reilly
- Google's Python class:
<https://www.youtube.com/watch?v=tKTZoB2Vjukxo>
- <https://docs.python.org/3.5/tutorial>
- *Lots of other materials available on the web*

Pandas example

+ Code + Text

The New York Times Covid-19 Database

The New York Times Covid-19 Database is a county-level database of confirmed cases and deaths, compiled from state and local governments and health departments across the United States. The initial release of the database was on Thursday, March 26, 2020, and it is updated daily.

These data have fueled many articles and graphics by The Times; these are updated regularly at <https://www.nytimes.com/interactive/2020/us/coronavirus-us-cases.html>. The Times has created many visualizations that are effective communications of important information about the pandemic.

The data are publically available via GitHub: <https://github.com/nytimes/covid-19-data>. In this illustration we will only use the data aggregated at the state level.

```
[ ] import pandas as pd
import numpy as np

%matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')

[ ] covid_table = pd.read_csv("https://raw.githubusercontent.com/nytimes/covid-19-data/master/us-states.csv")
covid_table = covid_table.drop('fips', axis=1)
covid_table.tail(20)
```

	date	state	cases	deaths
30464	2021-09-07	North Dakota	119995	1596
30465	2021-09-07	Northern Mariana Islands	248	2
30466	2021-09-07	Ohio	1262018	21020
30467	2021-09-07	Oklahoma	570923	8001
30468	2021-09-07	Oregon	200610	2202

Some Terminology

- supervised vs. unsupervised
- classification vs. regression
- prediction vs. inference

Supervised Learning vs. Unsupervised Learning

Supervised learning:

- Given a set of (x, y) , learn to predict y using x .
- e.g.
 - ▶ Predicting whether a loan will default based on customer characteristics

Supervised Learning vs. Unsupervised Learning

Supervised learning:

- Given a set of (x, y) , learn to predict y using x .
- e.g.
 - ▶ Predicting whether a loan will default based on customer characteristics

Unsupervised learning:

- Given a set of x , learn underlying structure or relationships of x .
- e.g.
 - ▶ Identifying market segments with similar spending patterns.

Classification vs. Regression

The Income dataset:

Education	Seniority	Income
21.58621	113.1034	99.91717
18.27586	119.3103	92.57913
12.06897	100.6897	34.67873
17.03448	187.5862	78.70281
19.93103	20.0000	68.00992
18.27586	26.2069	71.50449

Information for 30 *simulated individuals.*

Classification vs. Regression

The Income dataset:

Education	Seniority	Income
21.58621	113.1034	99.91717
18.27586	119.3103	92.57913
12.06897	100.6897	34.67873
17.03448	187.5862	78.70281
19.93103	20.0000	68.00992
18.27586	26.2069	71.50449

Regression: Model **income** based on other characteristics.

Information for 30 *simulated individuals*.

Classification vs. Regression

The Income dataset:

Education	Seniority	Income
21.58621	113.1034	99.91717
18.27586	119.3103	92.57913
12.06897	100.6897	34.67873
17.03448	187.5862	78.70281
19.93103	20.0000	68.00992
18.27586	26.2069	71.50449

Information for 30 *simulated individuals*.

Regression: Model **income** based on other characteristics.

Classification: Model **whether someone will earn above the median income** based on other characteristics.

Inference vs. Prediction

The Income dataset:

Education	Seniority	Income
21.58621	113.1034	99.91717
18.27586	119.3103	92.57913
12.06897	100.6897	34.67873
17.03448	187.5862	78.70281
19.93103	20.0000	68.00992
18.27586	26.2069	71.50449

Prediction: accurately predict Y for new observations

Information for 30 *simulated individuals.*

Inference vs. Prediction

The Income dataset:

Education	Seniority	Income
21.58621	113.1034	99.91717
18.27586	119.3103	92.57913
12.06897	100.6897	34.67873
17.03448	187.5862	78.70281
19.93103	20.0000	68.00992
18.27586	26.2069	71.50449

Information for 30 *simulated individuals.*

Prediction: accurately predict Y for new observations

Inference: explain the underlying relationship between Y and X

Example: Handwritten Digit Recognition

- Data: images of handwritten digits (grayscale pixel values)
- Classify images as digits 0 to 9.

0	4	1	9	2	1	3	1	4	3
5	3	6	1	7	2	8	6	9	4
0	9	1	1	2	4	3	2	7	3
8	6	9	0	5	6	0	7	6	1
8	7	9	3	9	8	5	9	3	3
0	7	4	9	8	0	9	4	1	4
4	6	0	4	5	6	1	0	0	1
7	1	6	3	0	2	1	1	7	9
0	2	6	7	8	3	9	0	4	6
7	4	6	8	0	7	8	3	1	5

Example: Handwritten Digit Recognition

0	4	1	9	2	1	3	1	4	3
5	3	6	1	7	2	8	6	9	4
0	9	1	1	2	4	3	2	7	3
8	6	9	0	5	6	0	7	6	1
8	7	9	3	9	8	5	9	3	3
0	7	4	9	8	0	9	4	1	4
4	6	0	4	5	6	1	0	0	1
7	1	6	3	0	2	1	1	7	9
0	2	6	7	8	3	9	0	4	6
7	4	6	8	0	7	8	3	1	5

- Data: images of handwritten digits (grayscale pixel values)
- Classify images as digits 0 to 9.

80322 - 4129 80206

40004 44310

37878 05453

~~35502~~ 75216

35460 44209

Summary

- Two cultures: model based and prediction based
- Python, pandas, and linear regression example with Covid-19 data

Next week: Linear regression and classification