S&DS 265 / 565
Introductory Machine Learning

Stochastic Gradient Descent

January 28

Yale

Reminders

- Quiz 1 will be live after class (2:30)
  - ▶ Taken online, on Canvas
  - ▶ 48 hour window
  - ▶ 20 minutes once started
  - ▶ Open notes, etc.
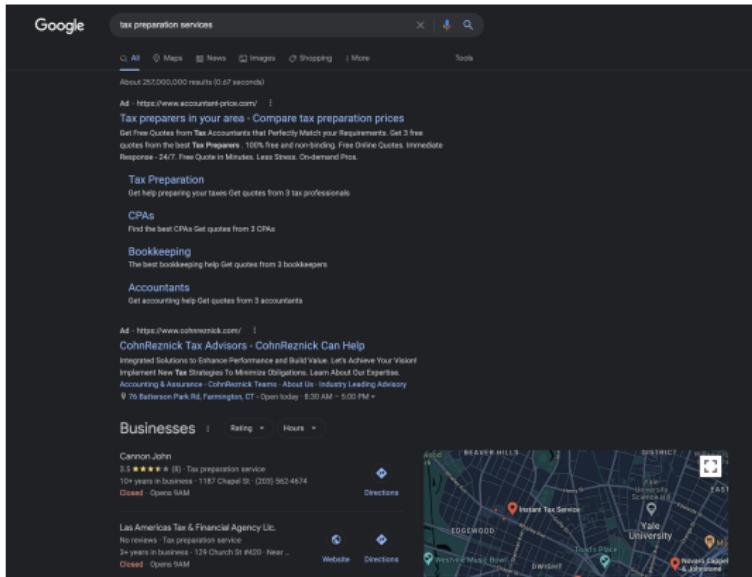- Assn 1 due week from tomorrow

Outline for today

- Continue discussion of generative/discriminative
- Stochastic gradient descent
- Application to logistic regression
- Regularization
- Learning rate and scaling
- Jupyter notebook example

Stochastic gradient descent

- Suppose that we want to fit a really big model, where the number of samples *n* and number of variables *p* are very large

- The classical algorithms in standard software packages will fail

- How can we train such models?

Recall: Classification tasks

- Ad click-through prediction. Predict whether or not a user will click on an ad presented. Used for ranking ads and setting prices.

Recall:Classification tasks

- Ad click-through prediction. Predict whether or not a user will click on an ad presented. Used for ranking ads and setting prices.

Example

- We want to classify ads according to whether or not they will be clicked on by a user

- We have a very large collection of training data

- Ads are represented in terms of a sparse list of features

  $1 \mid 5 : 1.1789641e{-}01 \quad 39 : 6.0373064e{-}02 \quad 45 : 1.3163488e{-}01$

- The dataset is too large to load into memory, and the number of features is also very large

- New data are continually arriving

- How can we efficiently train a classifier?

Online learning

We will introduce a method that

- Reads in the data points one (or a few) at a time
- Updates the model for each sample
- Exploits sparsity of the features
- Uses little memory, never reads in the entire dataset

Stochastic gradient descent

Initialize all parameters to zero: $\beta_j = 0$, $j = 1, \ldots, p$.

Read through the data one record at a time, and update the model.

1. Read data item x
2. Make a prediction $\widehat{y}(x)$
3. Observe the true response/label y
4. Update the parameters β so $\widehat{y}$ is closer to y

Stochastic gradient descent

To begin, suppose we are doing *linear regression*. We initialize all parameters to zero: $\beta_j = 0$, $j = 1, \ldots, p$.

We read through the data one record at a time, and update the model.

1. Read data item x
2. Make a prediction $\widehat{y}(x) = \sum_{j=1}^{p} \beta_j x_j$
3. Observe the true response/label y
4. Update the parameters β so $\widehat{y}$ is closer to y

SGD idea

Here's the idea:

- For each parameter β_j, see what happens to the loss if that parameter is increased a little bit.

- If the loss goes down (up), then increase (decrease) β_j proportionately

- Do this simultaneously for all of the parameters

- Rinse and repeat

SGD idea

Change β_j by a little bit:

$$\beta_j \to \beta_j + \varepsilon$$

What happens to the squared error?

$$\begin{aligned}
\tfrac{1}{2}(y - \widehat{y})^2 &\to \tfrac{1}{2}(y - \widehat{y} - \varepsilon x_j)^2 \\
&\approx \tfrac{1}{2}(y - \widehat{y})^2 - (y - \widehat{y})x_j\, \varepsilon \\
&= \tfrac{1}{2}(y - \widehat{y})^2 + \underbrace{(\widehat{y} - y)x_j}_{g}\, \varepsilon
\end{aligned}$$

SGD idea

We then change the parameter as follows:

$$\begin{aligned}
\beta_j &\to \beta_j - \eta g \\
&= \beta_j - \eta \underbrace{(\widehat{y} - y) x_j}_{g} \\
&= \beta_j + \eta (y - \widehat{y}) x_j
\end{aligned}$$

with $\eta > 0$ a small positive "step size".

SGD idea

Why is this a good idea? With this choice of $\varepsilon = -\eta g$ the squared error decreases:

$$\begin{aligned}
\tfrac{1}{2}(y - \widehat{y})^2 &\to \tfrac{1}{2}(y - \widehat{y} - \varepsilon x_j)^2 \\
&\approx \tfrac{1}{2}(y - \widehat{y})^2 + \varepsilon g \\
&= \tfrac{1}{2}(y - \widehat{y})^2 - \eta g^2 \\
&< \tfrac{1}{2}(y - \widehat{y})^2
\end{aligned}$$

so we're moving "downhill"!

SGD for general loss

Suppose $L(y, \beta^T x)$ is the loss for an input (x, y), e.g., $(y - \beta^T x)^2$

SGD update:

$$\beta_j \longleftarrow \beta_j - \eta \frac{\partial L(y, \boldsymbol{\beta}^T x)}{\partial \beta_j}$$

$$\boldsymbol{\beta} \longleftarrow \boldsymbol{\beta} - \eta \nabla_\beta L(y, \boldsymbol{\beta}^T x) \quad \text{(vector notation)}$$

- η is the *learning rate* or "step size"
- Needs to be chosen carefully, getting smaller over time

Gradient descent for general loss

If $L(\boldsymbol{\beta})$ is the loss function over subset of training set:
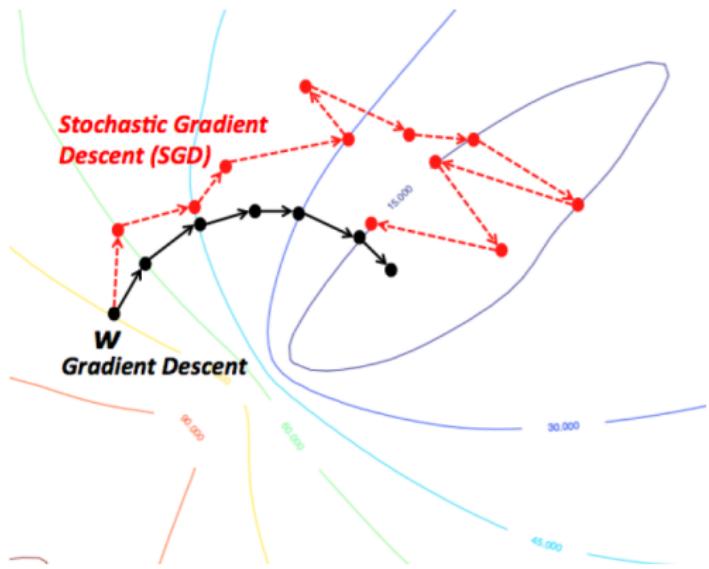
$$L(\boldsymbol{\beta} + \eta\boldsymbol{v}) \approx L(\boldsymbol{\beta}) + \eta\boldsymbol{v}^T\nabla L(\boldsymbol{\beta})$$
$$L(\boldsymbol{\beta} - \eta\nabla L(\boldsymbol{\beta})) \approx L(\boldsymbol{\beta}) - \eta\|\nabla L(\boldsymbol{\beta})\|^2$$

This is why gradient descent is going downhill — if η is small enough.

"Batch" gradient descent uses the entire training set in each step of gradient descent.

Stochastic gradient descent computes a quick approximation to this gradient, using only a single or a small "mini-batch" of data points

Batch vs. stochastic gradient descent

https://wikidocs.net/3413

SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose $y = 1$ and probability $p(x)$ is high?

SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose $y = 1$ and probability $p(x)$ is high? *small change*
- Suppose $y = 1$ and probability $p(x)$ is small?

SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose $y = 1$ and probability $p(x)$ is high? *small change*
- Suppose $y = 1$ and probability $p(x)$ is small? *big change* ↑
- Suppose $y = 0$ and probability $p(x)$ is small?

SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose $y = 1$ and probability $p(x)$ is high? *small change*
- Suppose $y = 1$ and probability $p(x)$ is small? *big change* $\uparrow$
- Suppose $y = 0$ and probability $p(x)$ is small? *small change*
- Suppose $y = 0$ and probability $p(x)$ is big?

SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - p(x))x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - p(x))x_j^2$$

$$p(x) = \frac{1}{1 + \exp(-\beta^T x)}$$

Case checking:

- Suppose $y = 1$ and probability $p(x)$ is high? *small change*
- Suppose $y = 1$ and probability $p(x)$ is small? *big change* $\uparrow$
- Suppose $y = 0$ and probability $p(x)$ is small? *small change*
- Suppose $y = 0$ and probability $p(x)$ is big? *big change* $\downarrow$

SGD: choice of step size

In theory, we need to let the step size η decrease as the algorithm progresses.

This prevents the estimates from oscillating back and forth without converging.

Demo

Open the demo notebook `sgd.ipynb` and follow along...

SGD: Scaling

We generally want to "standardize" each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves "looking ahead" to compute the mean and variance, and destroys the online property of the algorithm

SGD: Scaling

We generally want to "standardize" each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves "looking ahead" to compute the mean and variance, and destroys the online property of the algorithm

Solution: The mean and variance can be updated in an online manner, in constant time, by storing auxiliary variables for each component j.

SGD: Regularization

A "ridge" penalty $\frac{1}{2}\lambda \sum_{j=1}^{p} \beta_j^2$ is easily handled.

Gradient changes by an additive term $\lambda\beta$. Update becomes

$$\begin{aligned}
\beta_j &\longleftarrow & \beta_j + \eta\{(y - p(x))x_j - \lambda\beta_j\} \\
&=& (1 - \eta\lambda)\beta_j + \eta(y - p(x))x_j
\end{aligned}$$

Observe that this "does the right thing" whether β_j wants to be large positive or negative.

- *The penalty shrinks β_j toward zero*

What did we learn today?

- Stochastic gradient descent is a simple algorithm that can be applied to large classification and regression problems

- A parameter is updated according to how much the loss changes when that parameter is changed by a little bit

- This is the "go to" algorithm for fitting large or complex machine learning models

- Choosing the learning rate is a little tricky